

EARLY BIRD

AI Engineer | Python · PyTorch · LangChain · LLMs · RAG · MLOps · AWS
0300-1234567 · earlybird.ai@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | Lahore, Pakistan

PROFESSIONAL SUMMARY

Results-driven AI Engineer with hands-on experience designing, training, and deploying machine learning and large language model (LLM) systems in production. Proficient across the modern AI stack — Python, PyTorch, LangChain, and vector databases — with cloud deployment experience on AWS (EC2, S3, Lambda, SageMaker) and Hugging Face Spaces. Delivered measurable accuracy and latency improvements, retrieval-augmented generation (RAG) pipelines, and LLM evaluation frameworks in Agile team environments. Passionate about trustworthy AI, scalable ML architecture, and rigorous evaluation.

TECHNICAL SKILLS

Machine Learning & AI: Python, PyTorch, TensorFlow, Scikit-learn, Transformers, Fine-Tuning, Prompt Engineering, RAG, LangChain, LangSmith

NLP & LLM Tooling: OpenAI API, Grok API, Hugging Face, Vector Databases (Pinecone, Chroma, FAISS), Embeddings, Semantic Search

Backend & APIs: FastAPI, Flask, Node.js, REST API Design, WebSockets, JWT Authentication

Data & MLOps: Pandas, NumPy, Docker, MLflow, CI/CD Pipelines, Model Monitoring, Airflow, Experiment Tracking

Cloud & DevOps: AWS (EC2, S3, Lambda, SageMaker), Hugging Face Spaces, Git, GitHub Actions, Nginx

Databases: PostgreSQL, MongoDB, Redis, Vector Stores, Query Optimisation

WORK EXPERIENCE

AI Engineer | Nimbus AI Labs

Sep 2025 – date | Lahore, PK

- Designed and deployed a retrieval-augmented generation (RAG) pipeline using LangChain and a vector database, reducing hallucinated responses by 35% and improving answer relevance across a 10,000+ document knowledge base.
- Built and fine-tuned transformer-based NLP models for text classification and entity extraction, improving downstream accuracy by 22% over the previous rule-based system.
- Implemented LangSmith-based evaluation pipelines to track model accuracy, latency, and cost across releases, and integrated evaluation gates into CI/CD so regressions were caught before deployment.
- Collaborated in Agile sprints with product and backend teams; maintained 100% on-time sprint delivery across all release cycles.

Junior Machine Learning Engineer | Vertex Analytics

May 2024 – Sep 2024 | Lahore, PK

- Shipped 4+ ML-powered features end-to-end: data preprocessing pipelines, model training in PyTorch/Scikit-learn, and FastAPI endpoints for real-time inference.
- Optimised model inference latency by 30% through batching, caching, and quantisation, enabling deployment on cost-efficient AWS instances.
- Integrated third-party LLM APIs (OpenAI, Hugging Face Inference) into internal tools, boosting analyst productivity by 18%.

PROJECTS

TRUTH-CHECKED RESUME TAILOR — Fabrication-Guarded AI Resume Assistant (LangChain + LangSmith) [GitHub](#)

- Built an AI pipeline that tailors resumes to job postings while guaranteeing zero fabricated claims, using a fact-extraction step, a vector database for meaning-based matching, and a second-pass LLM verifier.
- Tracked fabrication rate, requirement coverage, and cost via LangSmith across a 45-case evaluation set; deployed the app on Hugging Face Spaces with a GitHub Actions gate blocking releases on regression.

DOCUCCHAT — RAG-Based Document Q&A Assistant (Full Stack) [GitHub](#)

- Built a retrieval-augmented chatbot over user-uploaded PDFs using LangChain, FAISS embeddings, and OpenAI's API, with a React frontend and FastAPI backend.
- Reduced average response latency by 25% through embedding caching and asynchronous request handling; deployed on AWS EC2 behind Nginx.

CHURNSENSE — Customer Churn Prediction Platform (ML + MLOps) [GitHub](#)

- Trained and deployed a gradient-boosted churn prediction model achieving 89% AUC, with automated retraining and drift monitoring via MLflow.
- Containerised the training and inference services with Docker and set up a GitHub Actions CI/CD pipeline for automated testing and deployment.

EDUCATION

BS Computer Science — COMSATS University Islamabad, Lahore Campus CGPA: 3.45 / 4.0

Relevant Coursework: Machine Learning · Deep Learning · Natural Language Processing · Data Structures & Algorithms · Database Systems · Distributed Systems · Artificial Intelligence · Software Design & Architecture

CERTIFICATIONS

- Deep Learning Specialization · DeepLearning.AI (Coursera)
- LangChain for LLM Application Development · DeepLearning.AI
- AWS Certified Machine Learning – Specialty · Amazon Web Services